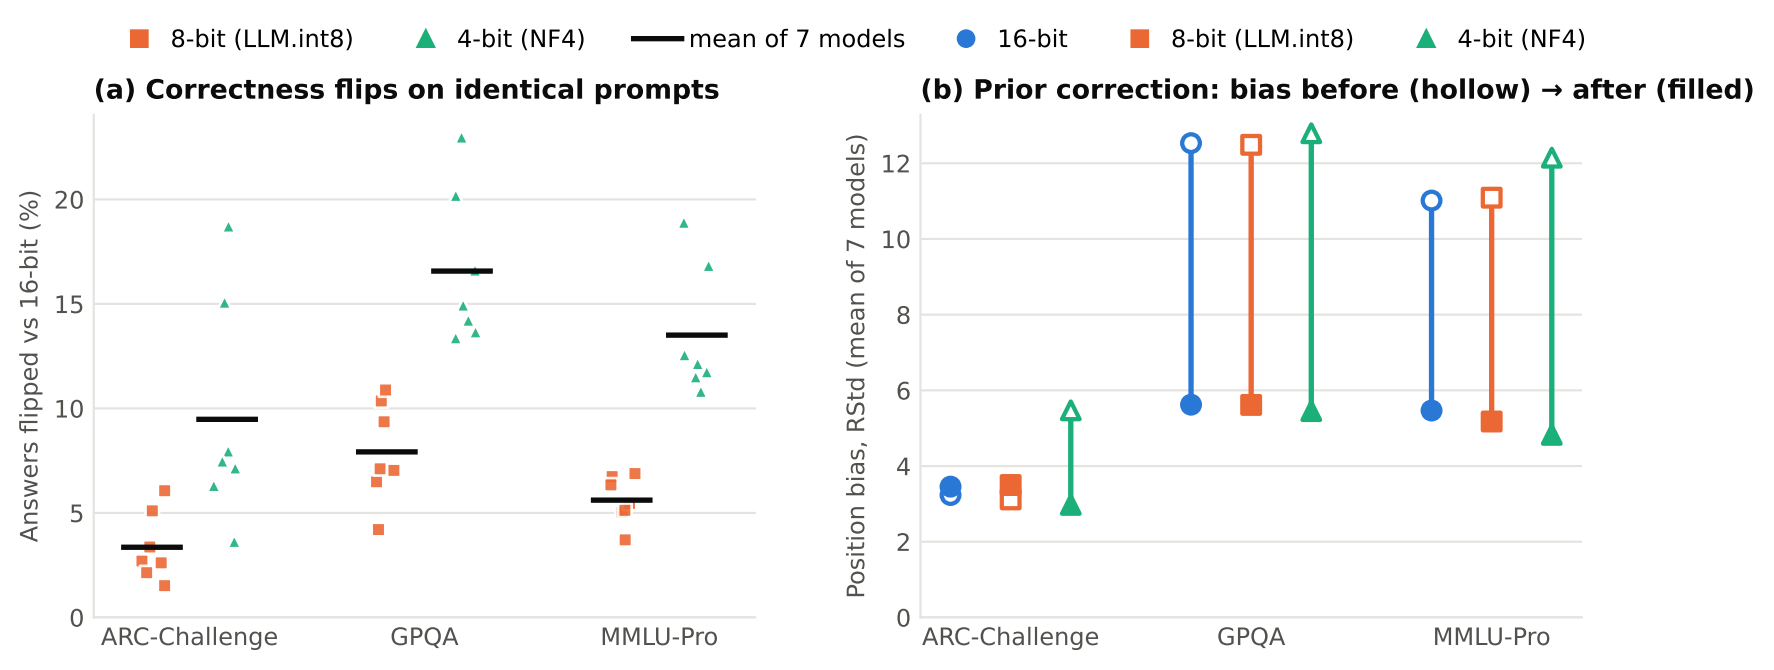
8-bit (LLM.int8)
4-bit (NF4)
mean of 7 models
16-bit
8-bit (LLM.int8)
4-bit (NF4)
(a) Correctness flips on identical prompts
Answers flipped vs 16-bit (%)
0
5
10
15
20
ARC-Challenge
GPQA
MMLU-Pro
(b) Prior correction: bias before (hollow) → after (filled)
Position bias, RStd (mean of 7 models)
0
2
4
6
8
10
12
ARC-Challenge
GPQA
MMLU-Pro